

# P6 – Scientific Programming

Marcus Mohr  
Jens Oeser

Geophysics Section  
Department of Earth and Environmental Sciences  
Ludwig-Maximilians-Universität München

SoSe 2021

## Part #8

# Operators: Bits & Bobs

'casts', sizeof, ternary if  
alignment & padding

## Casts (1/2)

- We had discussed that in expressions involving operands of mixed types an **implicit type conversion** will occur.
- This happens by “promoting” or “upcasting” operands to “higher” types.
- We can explicitly express this by performing an **explicit type conversion**.
- Uses a **cast** operator (datatype).
- Only works with basic datatypes and pointers.

```
int i = 2;  
double x = 5.0, y;  
  
// implicit type conversion  
y = i * x;
```

```
// explicit type conversion  
y = (double)i * x;
```

```
signed char c = 5;  
int m = (int)c;  
  
double x = 1.2;  
float f = (float)x;
```

## Casts (2/2)

- We can also **downcast** variables and constants:
  - ▶ Introduces a potential data range issue.  
conversion to 'signed char' from 'int' may alter its value [-Wconversion] `signed char s = m;`
  - ▶ Downcasting a floating-point value to an integer involves **truncation**.
- What will the code on the right print?

```
int m = 42;
signed char s = m;
unsigned char u =
    (unsigned char)m;

double x = -1.97;
s = (signed char)x;
u = (unsigned char)x;

printf( "x = %f ", x );
printf( "s = %d ", s );
printf( "u = %u\n", u );
```

## Casts (2/2)

- We can also **downcast** variables and constants:
  - ▶ Introduces a potential data range issue.  
conversion to 'signed char' from 'int' may alter its value [-Wconversion] `signed char s = m;`
  - ▶ Downcasting a floating-point value to an integer involves **truncation**.
- What will the code on the right print?

`x = -1.97, s = -1, u = 255`

```
int m = 42;
signed char s = m;
unsigned char u =
    (unsigned char)m;

double x = -1.97;
s = (signed char)x;
u = (unsigned char)x;

printf( "x = %f ", x );
printf( "s = %d ", s );
printf( "u = %u\n", u );
```

## Example: Printing Booleans

- C does not provide a conversion specifier for data of type `_Bool`.

```
bool logical = true;

if( logical ) {
    printf( "logical is true\n" );
}
else {
    printf( "logical is false\n" );
}
```

- There is a more compact alternative:

```
printf( "logical is %s\n",
        logical ? "true": "false" );
```

# Conditional Operator (1/2)

- `?:` is called “conditional operator”, “inline if” or “ternary operator”.
- It consists of three parts:

```
<expression A> ? <expression B> : <expression C>
```

- and works as follows:
  - ① `<expression A>` is evaluated
  - ② if result is true, `<expression B>` is evaluated and its result returned
  - ③ if result is false, `<expression C>` is evaluated and its result returned

## Conditional Operator (2/2)

```
<expression A> ? <expression B> : <expression C>
```

- The return type of the ternary operator is the higher one of those of the expressions B and C

```
k > 0 ? 10 : -2.0
```

always returns a double.

- The return type can also be ignored like here

```
x > 0.0 ? printf( "positive" ) : printf( "not positive" );
```



## Operator: sizeof (1/2)

- The `sizeof()` operator allows to query the memory size in bytes of a
  - ▶ datatype
  - ▶ variable
  - ▶ literal
- Its return value is of type `size_t`. This is a compiler specific unsigned integer type.

```
printf( "[char ] %lu bytes, ", sizeof(char) );  
printf( "[short] %lu bytes, ", sizeof(short) );  
printf( "[int   ] %lu bytes\n", sizeof(int) );
```

prints: [char ] 1 bytes, [short] 2 bytes, [int ] 4 bytes

## Operator: sizeof (1/2)

- The `sizeof()` operator allows to query the memory size in bytes of a
  - ▶ datatype
  - ▶ variable
  - ▶ literal
- Its return value is of type `size_t`. This is a compiler specific unsigned integer type.

```
float f;
printf( "[float] %lu bytes, ", sizeof(float) );
printf( "[f      ] %lu bytes, ", sizeof(f)      );
printf( "[3.2f ] %lu bytes\n", sizeof(3.2f) );
```

prints: [float] 4 bytes, [f ] 4 bytes, [3.2f ] 4 bytes

## Operator: sizeof (3/3)

Let's examine more literals:

```
printf( "[true          ] %lu bytes\n", sizeof(true)          );  
printf( "[(bool>true)] %lu bytes\n", sizeof((bool>true) ) );  
printf( "['A'          ] %lu bytes\n", sizeof('A')            );  
printf( "[(char)'A'   ] %lu bytes\n", sizeof((char)'A')      );  
printf( "[\"my size?\" ] %lu bytes", sizeof("my size?") );
```

Code gives us:

[true          ] 4 bytes	← true internally is an integer literal
[(bool>true)] 1 bytes	
['A'          ] 4 bytes	← letters internally are integer literals
[(char)'A'   ] 1 bytes	
["my size?" ] 9 bytes	← string terminated by trailing '\0'

# Alignment and Padding (1/8)

- We can apply `sizeof()` also to derived datatypes:

```
typedef struct {  
    int iVal;  
    float fVal;  
} compound1;  
  
printf( "[compound1] %lu bytes\n", sizeof(compound1) );
```

- The above results in `[compound1] 8 bytes`.
- That's the sum of `[int] 4 bytes` and `[float] 4 bytes`.

## Alignment and Padding (2/8)

- Let's change the float to a double

```
typedef struct {  
    int iVal;  
    double dVal;  
} compound2;  
  
printf( "[compound2] %lu bytes\n", sizeof(compound2) );
```

- The above results in **[compound2] 16 bytes**.
- But that is larger than: **[int] 4 bytes + [double] 8 bytes?**

## Alignment and Padding (3/8)

```
58 typedef struct {  
59     short s1;  
60     double dVal;  
61     short s2;  
62 } compound5;  
63  
64 printf( "[compound5] %lu bytes\n", sizeof(compound5) );
```

- We might expect  $(2 + 8 + 2) = 12$  bytes as answer, but get **24 bytes!**
- What happens here is called **padding** to ensure correct **alignment**.

## Alignment and Padding (4/8)

- Storing the three components consecutively one after another requires 12 bytes:




- However, access to dVal would be unaligned.
- Assuming a 64-bit system, a word occupies 8 bytes. Memory access (read/write) always transfers a collection of complete words.
- To ensure that dVal can be loaded/stored with a single access it must be aligned to word boundaries, i.e. multiples of 8 bytes.

## Alignment and Padding (5/8)

- Compiler enforces alignment of dVal by “padding” our struct with 6 meaningless bytes:



- To make the whole struct align it also adds 6 meaningless bytes at the end:





## Alignment and Padding (6/8)

- Adding compiler option `-Wpadded` to GCC gives this report:

```
warning: padding struct to align 'dVal' [-Wpadded]
    double dVal;
        ~~~~

warning: padding struct size to alignment boundary [-Wpadded]
    } compound5;
    ^
```

## Alignment and Padding (7/8)

- C standards leaves details of alignment to the implementation.
- Different compilers might handle it differently. Clang does the same:

```
warning: padding struct 'compound5' with 6 bytes to align  
      'dVal' [-Wpadded]
```

```
      double dVal;  
      ^
```

```
warning: padding size of 'compound5' with 6 bytes to alignment  
      boundary [-Wpadded]
```

```
      typedef struct {  
      ^
```

## Alignment and Padding (8/8)

- Reordering the components in our struct reduces the amount of padding:  
[C compiler is not allowed to do this!]

```
typedef struct {  
    double dVal;  
    short s1;  
    short s2;  
} compound6;
```

- This now only requires 16 bytes:

