

P6 – Scientific Programming

Marcus Mohr
Jens Oeser

Geophysics Section
Department of Earth and Environmental Sciences
Ludwig-Maximilians-Universität München

SoSe 2021

Part #11

Pointers

memory address, dereferencing, pointer arithmetic,
dynamic memory management

Variables (again)


Aspects

In an imperative language four aspects make up a variable:

- ① the memory area serving as storage container for the data stored in the variable
- ② the address of the memory area
- ③ the data stored in the memory area
- ④ the variable's name to access the data in the memory area

What is a Pointer?

- Computer memory can be imagined as a sequence of drawers, that store your data



- Pointers are **variables for storing memory addresses** (drawer numbers).
- They allow us to locate our data in memory.

Pointer Definition

- to declare a pointer variable one uses the * token
- we note that a pointer variable always is associated with a **data type**, i.e. it can only store addresses for objects of that fixed type
- thus, `iPtr` is called a **pointer to int** or briefly an **int pointer**; and so on
- spaces around the * do not matter; though `int* iPtr` more clearly expresses that `iPtr` is an **int pointer**

```
int* iPtr;  
double *dPtr;  
float * fPtr;
```

```
// careful with lists:  
  
short s, *sPtr;  
// s (short)  
// sPtr (short pointer)  
  
int* iPtr1, iPtr2;  
// iPtr (int pointer)  
// iPtr2 (int)
```

Obtaining & Storing Addresses

- In order to assign a value to a pointer, we need a way to find the memory address of an object.
- For this there is the **address operator &**.
- We can also copy the value from one pointer into another one.

```
double dVal;  
double *dPtr = &dVal;  
int *iPtr1, *iPtr2;  
int iVal;  
iPtr1 = &iVal;  
  
iPtr2 = iPtr1;  
iPtr2 = dPtr;    // dubious as pointers have different type!
```

- After the assignment **iPtr1 = &iVal;** we say that **iPtr1 points to iVal**.

Addresses (1/2)

- Most data types occupy more than a single word. So what is the address of a variable?
- Let us examine the following example:

```
typedef struct { float fVal; int iVal; } compound;  
compound var;
```

- In memory var will (typically) occupy eight bytes



Addresses (2/2)

- this code

```
compound* cPtr = &var;  
float* fPtr = &var.fVal;  
int* iPtr = &var.iVal;  
printf( "cPtr = %p\n", cPtr );  
printf( "fPtr = %p\n", fPtr );  
printf( "iPtr = %p\n", iPtr );
```

- might print

```
cPtr = 0x7ffc23fc2bb0  
fPtr = 0x7ffc23fc2bb0  
iPtr = 0x7ffc23fc2bb4
```

(address details are runtime dependent;
distance between fPtr and iPtr is not)



- an object's address is always that of its first byte

Dereferencing

- The ***** operator allows to access the memory location a pointer points to.
- This is called **dereferencing** the pointer.

```
int a = 3, b = 0;  
int* ptr;  
ptr = &a;  
b = *ptr;  
*ptr = 5;  
printf( "a=%d, b=%d\n", a, b );
```

- Dereferencing allows to read or write data from that memory address.
- Code above will print **a=5, b=3** .

Null Pointer

- What will happen, when we run this code?

```
int* ptr;  
*ptr = 5;
```

Null Pointer

- What will happen, when we run this code?

```
int* ptr;  
*ptr = 5;
```

- Most likely this:

```
Segmentation fault (core dumped)
```

The pointer variable `ptr` was not initialised. Thus, it will point to some arbitrary memory address!

Null Pointer

- What will happen, when we run this code?

```
int* ptr;  
*ptr = 5;
```

- Most likely this:

```
Segmentation fault (core dumped)
```

The pointer variable `ptr` was not initialised. Thus, it will point to some arbitrary memory address!

- How can we find out if a pointer is unitialised?

Null Pointer

- What will happen, when we run this code?

```
int* ptr;  
*ptr = 5;
```

- Most likely this:

```
Segmentation fault (core dumped)
```

The pointer variable `ptr` was not initialised. Thus, it will point to some arbitrary memory address!


- How can we find out if a pointer is unitialised?
- We initialise it to an invalid, but well defined value. C provides the macro `NULL` for this purpose:

```
double* dPtr = NULL; // good style
```

Pointer Arithmetics

```
short int a, b, c, *ptr;
```

Assume that a short integer occupies two drawers and that a, b and c are stored one after the other



```
a
ptr = &a;
ptr = ptr + 2;
a = *ptr;
```

```
b
/* ptr -> drawer #20 */
/* ptr -> drawer #24 */
/* a = c */
```

Example: Endianness Test (1/2)

```
#include <stdio.h>

// for uint*_t types
#include <stdint.h>

int main() {

    uint16_t i = 1u;
    uint8_t* c = (uint8_t*)&i;


    if( *c )
        printf( "little endian\n" );
    else
        printf( "big endian\n" );

}
```

- Code on the left will report the machine's endianness.
- Works as follows:
 - ➊ define a 2-byte unsigned integer variable and initialise it to 1
 - ➋ get its (starting) address and cast it to a 1-byte unsigned integer pointer
 - ➌ dereference pointer and check value


Example: Endianness Test (2/2)

little endian:



code accesses value
stored in first byte;

big endian:



if value of bit pattern
in non-zero we have a
little endian system

Example: Taming the Panda

```
#include <string.h>
#include <stdio.h>

int main() {

    char v1[] = "Eats, shoots and leaves";
    //          01234567890123456789012
    //          ^           ^
    char v2[ strlen( v1 ) ];

    printf( "'%s'\n", v1 );    // prints 'Eats, shoots and leaves'!

    strncpy( v2, v1, 4 );
    strcpy( &v2[4], &v1[5] );
    // or: strcpy( &v2[0]+4, &v1[0]+5 );
    // or: strcpy( v2+4, v1+5 );

    printf( "'%s'\n", v2 );    // prints 'Eats shoots and leaves'!
}
```

Another Example

Examine the following codelet:

```




1  int a;
2  int* x;
3  int** y;
4  y = &x;
5  x = &a;
6  **y = 1;
7  *x = a + **y;
8  a = *x + **y;


```

We assume that an int occupies 4 bytes and a pointer 8 bytes ...

... and that the variables are residing in memory one after the other starting at address 1,000.

memory occupation: **after line #3**

int	**	y		1012
int	*	x		1004
int		a		1000

 = garbage

Another Example

Examine the following codelet:

```

1  int a;
2  int* x;
3  int** y;
4  y = &x;
5  x = &a;
6  **y = 1;
7  *x = a + **y;
8  a = *x + **y;

```

We assume that an int occupies 4 bytes and a pointer 8 bytes ...

... and that the variables are residing in memory one after the other starting at address 1,000.

memory occupation: **after line #4**

int	**	y	1004	1016
int	*	x	⊗	1004
int		a	⊗	1000

⊗ = garbage

Another Example

Examine the following codelet:

```

1  int a;
2  int* x;
3  int** y;
4  y = &x;
5  x = &a;
6  **y = 1;
7  *x = a + **y;
8  a = *x + **y;

```

We assume that an int occupies 4 bytes and a pointer 8 bytes ...

... and that the variables are residing in memory one after the other starting at address 1,000.

memory occupation: **after line #5**

int	**	y	1004	1012
int	*	x	1000	1004
int		a	⊗	1000

⊗ = garbage

Another Example

Examine the following codelet:

```

1  int a;
2  int* x;
3  int** y;
4  y = &x;
5  x = &a;
6  **y = 1;
7  *x = a + **y;
8  a = *x + **y;

```

We assume that an int occupies 4 bytes and a pointer 8 bytes ...

... and that the variables are residing in memory one after the other starting at address 1,000.

memory occupation: **after line #6**

int	**	y	1004	1012
int	*	x	1000	1004
int		a	1	1000

⊗ = garbage

Another Example

Examine the following codelet:

```

1  int a;
2  int* x;
3  int** y;
4  y = &x;
5  x = &a;
6  **y = 1;
7  *x = a + **y;
8  a = *x + **y;
```

We assume that an int occupies 4 bytes and a pointer 8 bytes ...

... and that the variables are residing in memory one after the other starting at address 1,000.

memory occupation: **after line #7**

int	**	y	1004	1012
int	*	x	1000	1004
int		a	2	1000

⊗ = garbage

Another Example

Examine the following codelet:

```

1  int a;
2  int* x;
3  int** y;
4  y = &x;
5  x = &a;
6  **y = 1;
7  *x = a + **y;
8  a = *x + **y;

```

We assume that an int occupies 4 bytes and a pointer 8 bytes ...

... and that the variables are residing in memory one after the other starting at address 1,000.

memory occupation: **after line #8**

int	**	y	1004	1012
int	*	x	1000	1004
int		a	4	1000

⊗ = garbage

Member Access “Syntactic Sugar”

```
typedef struct {
    double x;
    int m;
} compound;

int main( void ) {

    compound cVar;
    cVar.x = 0.0;
    cVar.m = 1;

    compound* cPtr = &cVar;
    (*cPtr).m = 3;
    // *cPtr.m = 3;  <- error!
    cPtr->x = 2.0;
}
```


What are Pointers Good for?

- Causing confusion ;-)
- Dynamic memory management
- Changing variables from within subprograms
- Linked lists, tree data structures, ...

Dynamic Memory Management

- In many applications the amount of memory required is not known a priori at **compile time** (e.g. importing an unstructured mesh)
- Dynamic memory management means that memory is requested by the program at **run time** (from the OS).
- Pointers are closely related to dynamic memory management.
- Routines that allocate memory return **pointers** to the starting address of the allocated memory block.

Library Functions from `stdlib.h`

```
void* malloc(size_t size);
```

`malloc()` allocates `size` bytes and returns a pointer to the allocated memory. The memory is not cleared!

```
void* calloc(size_t nmemb, size_t size);
```

`calloc()` allocates memory for an array of `nmemb` elements of `size` bytes each and returns a pointer to the allocated memory. The memory is set to zero!

```
void free(void* ptr);
```

`free` frees the memory space pointed to by `ptr`, which must have been returned by a previous call to `malloc()` or `calloc()`.

Datatype void

- C has a special **void** datatype.
- It is **incomplete**, i.e. we cannot declare variables of type void.
- However, we can have void pointers
 - ▶ no pointer arithmetic![‡]
 - ▶ no dereferencing!
- We will encounter void again in the unit on functions.

[‡] GCC will allow this as language extension

```
#include <stdlib.h>

int main() {

    // not allowed!
    // void freeVar;

    void* freePtr;

    double x = 2.0;
    freePtr = (void*) &x;

    double y;

    // not allowed!
    // y = *freePtr;

    y = *(double*)freePtr;

}
```

Example: Allocate 1D Array (1/2)

Allocate memory for a vector of four doubles $(v_1, \dots, v_4)^T$:

```
double *v = NULL;  
v = (double*) calloc( 4, sizeof(double) );
```

Memory:

v_1	v_2	v_3	v_4
-------	-------	-------	-------

Access:

$*(v+0)$ $*(v+1)$ $*(v+2)$ $*(v+3)$

Shorter:

$v[0]$ $v[1]$ $v[2]$ $v[3]$

equivalent:

$$*(v + i) = v[i]$$

Example: Allocate 1D Array (2/2)

- With dynamic arrays we are not fixed to starting arrays at index 0!

```
double *v = NULL;  
v = (double*) calloc( 4, sizeof(double) );  
v--;      /* equiv to v = v - 1; */
```

- Now `v` points to a position `sizeof(double)` bytes in front of our allocated memory block!
- `v[1]` now points to v_1 (by pointer arithmetics).
- Technique is (nearly) legal, but **never** try to access `v[0]`
(But: out-of-bounds accesses are always a bad idea!)

```
#include<stdio.h>  // by Chr. Feichtinger
#include<stdlib.h>

int main() {
    int i, n;
    char* buffer = NULL;

    printf( "How long should the string be?\n" );
    scanf( "%d", &i );


    buffer = (char*)malloc( sizeof(char) * (i+1) );
    if( buffer == NULL ) {
        fprintf( stderr, "Memory allocation failed!\n" );
        exit(1);
    }
    for( n=0; n<i; ++n )
        buffer[n] = rand()%26 + 'a';

    buffer[i] = '\0';
    printf( "Random string: %s\n", buffer );

    free( buffer ); return 0; }
```


Dynamic Allocation of 2D Arrays

- Let u be a function “living” on a 2D grid: $u_{ij} := u(x_i, y_j)$



- We want to **dynamically allocate memory** to store the values of u .
- There are several ways to do this. We will examine three of them.

Variant #1



```
// u is of type pointer to pointer
// to double
double** u = NULL;
```

```
// allocate memory for holding
// addresses of mesh columns
u = (double**)
    malloc( nx * sizeof(double*) );
```

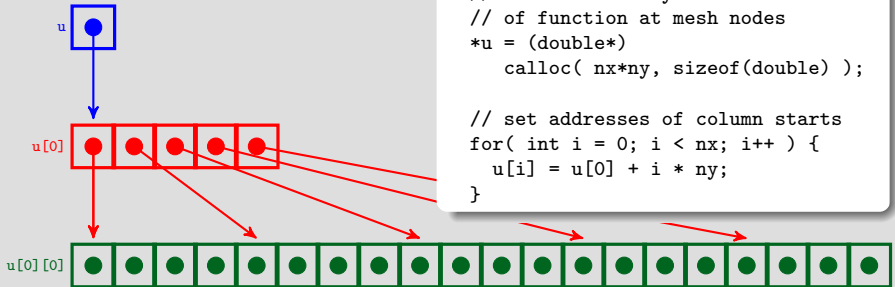
```
// allocate memory for entries
// in mesh columns
for( int i = 0; i < nx; i++ ) {
    u[i] = (double*)
        calloc( ny, sizeof(double) );
}
```

We can access values as: $u_{i,j} \equiv u[i][j] = *((u+i)+j)$

Variant #2

same as variant #1, but values
of u are stored in a contiguous
memory area

access again via $u[i][j]$



```
// u is of type pointer to pointer
// to double
double** u = NULL;

// allocate memory for holding
// addresses of columns starts
u = (double**)
    malloc( nx * sizeof(double*) );


// allocate memory for values
// of function at mesh nodes
*u = (double*)
    calloc( nx*ny, sizeof(double) );

// set addresses of column starts
for( int i = 0; i < nx; i++ ) {
    u[i] = u[0] + i * ny;
}
```

Variant #3

```
// u now is of type pointer to double  
double* u = NULL;
```

```
// we allocate one single chunk for  
// the function values at mesh nodes  
u = (double*)  
    calloc( nx*ny, sizeof(double) );
```



- similar to variant #2 we store all values of the function u consecutively in a contiguous memory area
- however, we skip the intermediate layer with the addresses of the column starting positions
- need to do address computation ourselves: $u_{i,j} \equiv u[i*ny+j]$


Comparison

- Pointer/address calculations for variants (1) and (2) are basically as expensive as the mapping in (3).
- Type (1) and (2) require additional storage for the pointer vectors.
(1) also demands more internal administration work for malloc than (2)
- Number of memory accesses to get value of $u_{i,j}$
 - ▶ Type (1): three memory accesses
 - ▶ Type (2): three memory accesses
 - ▶ Type (3): two memory accesses
- With (1) there is no control, where the data are placed in memory (paging and caching effects!).

Variant (2) seems a good compromise between efficiency and usability.

Adding Offsets

As with the 1D array allocation, we can also add offsets to our pointers to e.g. get one-based indexing for a matrix:



sizeof: static vs. dynamic

```
#include <stdlib.h>
#include <stdio.h>

int main( void ) {

    int staticVec[10] = {0};
    int* dynamicVec = (int*)calloc( 10, sizeof(int) );

    printf( "sizeof(staticVec)  = %2lu\n", sizeof(staticVec) );
    printf( "sizeof(dynamicVec) = %2lu\n", sizeof(dynamicVec) );

}
```

sizeof: static vs. dynamic

```
#include <stdlib.h>
#include <stdio.h>

int main( void ) {

    int staticVec[10] = {0};
    int* dynamicVec = (int*)calloc( 10, sizeof(int) );

    printf( "sizeof(staticVec)  = %2lu\n", sizeof(staticVec) );
    printf( "sizeof(dynamicVec) = %2lu\n", sizeof(dynamicVec) );

}
```

```
sizeof(staticVec) = 40
sizeof(dynamicVec) = 8
```

sizeof: static vs. dynamic

```
#include <stdlib.h>
#include <stdio.h>

int main( void ) {

    int staticVec[10] = {0};
    int* dynamicVec = (int*)calloc( 10, sizeof(int) );

    printf( "sizeof(staticVec)  = %2lu\n", sizeof(staticVec) );
    printf( "sizeof(dynamicVec) = %2lu\n", sizeof(dynamicVec) );

}
```

```
sizeof(staticVec) = 40
sizeof(dynamicVec) = 8    <-- size of a pointer variable!
```