

# P6 – Scientific Programming

Marcus Mohr

Jens Oeser

Geophysics Section  
Department of Earth and Environmental Sciences  
Ludwig-Maximilians-Universität München

SoSe 2021

## Part #15

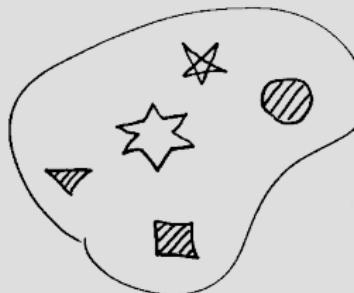
# Data Structures

abstract data-types, trees, stacks, queues,  
linked lists and containers

# Data Structures

we all know / have encountered data-structures

- 8 x Brezen
- 12 x Weißwürscht
- 6 x Weizen
- 1 x süßer Senf



## a shopping list

a set

a data-base record

**Marcus Mohr**  
(Numerical Simulation  
Department of Earth and  
Geophysics  
Munich University  
Theresienstr. 41  
80333 Munich  
Germany

Room: 411  
Phone: +49 (89) 2180-4230  
Fax: +49 (89) 2180-4205

E-Mail: marcus.mohr@geophysik.uni-muenchen.de  
Internet: <http://www.geophysik.uni-muenchen.de/~mohr>

## Data Structures (cont.)

## Definition:

An organization of information, usually in memory, for better algorithm efficiency, such as queue, stack, linked list, heap, dictionary, and tree.

or conceptual unity, such as the name and address of a person. It may include redundant information, such as length of the list or number of nodes in a subtree.

From: Dictionary of Algorithms and Data Structures

## Data Structures (cont.)

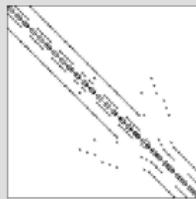
The main motivations for investigating/designing data structures are

- to find optimal data layouts
  - that allow to efficiently manipulate/administrate data
  - for certain applications
  - to design ways to handle associated data concertenly

Object-oriented programming takes this to a limit by encapsulating data and methods for manipulating them into objects/classes.

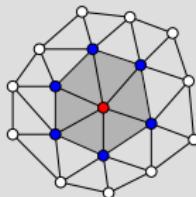
# Scientific Programming

Specialised data structures are not only the realm of computer science, but also valuable for applications in computational science.



Sparse matrices require specialised storage schemes, like e.g. [compressed row storage](#)

Complex simulations require a significant amount of steering parameters. We may want to group these into a single data structure for handling.



Efficient book-keeping on neighbouring elements for a FEM grid is non-trivial.

# Language Support

If one considers the concrete implementation of a data-structure in a certain programming language, one often speaks of an **abstract data-type** as opposed to the standard intrinsic types (integer, real, ...).

All **modern languages** (e.g. C, C++, Fortran90/95/2003) support the construction of abstract data-types using standard types as building blocks. These are often (e.g. in Fortran) denoted as **derived types**.

FORTRAN77, however, does not directly support the construction of derived types. (**too old**)

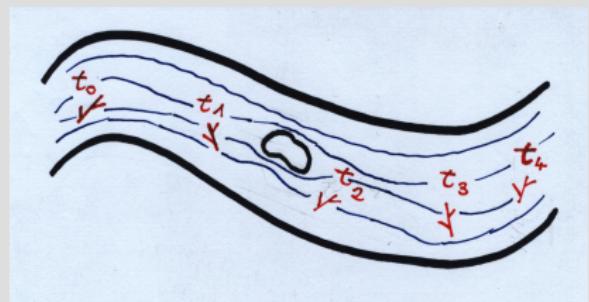
# Principle Concepts Overview

type	comment	access
array	fixed number of data-items of same type are stored contiguously, supported directly in most languages	by index
list	dynamic storage of arbitrarily many data-items, each item contains a link to its successor	traversal
stack	dynamic storage of arbitrarily many data-items, last item stored must be accessed first (operations: push & pop)	LIFO
queue	dynamic storage of arbitrarily many data-items, first item stored must be accessed first (operations: enqueue & dequeue)	FIFO
tree	special structure for modelling hierarchies, many specialisations like e.g. balanced trees, octrees, ...	
heap	implements set structure, elements associated with keys for ordering (→ priority queues)	

# Application: Particle Tracing

One possibility to visualise the behaviour of a flow pattern (used e.g. in computational fluid dynamics) is **particle tracing**

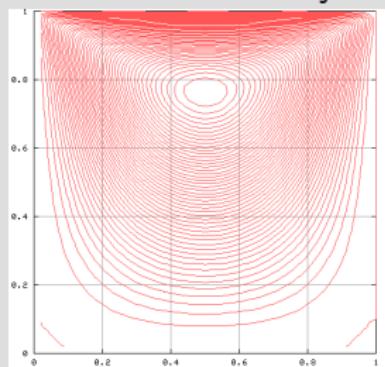
- at starting time insert particles into flow
- in each time step  $t_k$  compute current position of particles
- in intervals store current positions and visualise



“throw a stick into the river and watch it float”

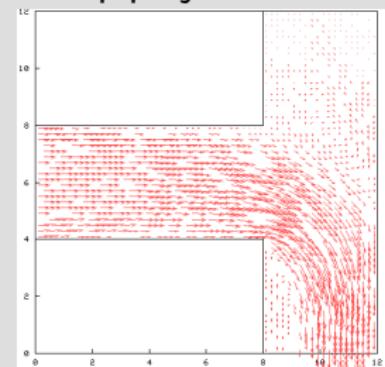
# Application: Particle Tracing (cont.)

lid driven cavity



simple case:  
number of particles  
is constant

pipe junction



more complex:  
number of particles  
changes over time

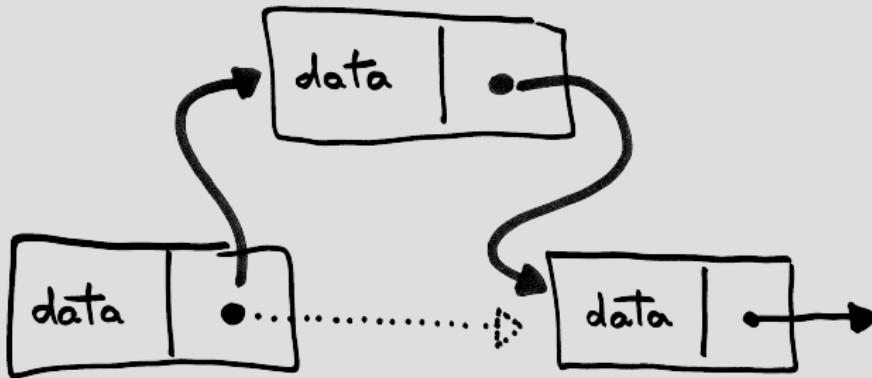
# Singly Linked Lists

- A suitable data structure for storing a changing number of data-items with repeated insertions and deletions is a singly linked list .
- The costs for inserting and deleting one data item are constant.



- Simplest form of a list. Each item contains a link to its successor.
- Last item needs a special NIL link to indicate end of list.

## Linked List: Insertion



- split list at desired insertion point (after left neighbour of new item)
  - insert link in left item to **new item**
  - insert link in **new item** to right item

# Linked List: Deletion



- split list in front of item to be deleted
- let predecessor of item point to its successor
- destroy item (if applicable)

# Linked Lists in C

- In a language with dynamic memory allocation (such as C) we can in principle store an arbitrarily long list (within the limits of the available memory)
- A link in this case is simple the memory address of the next list item (stored in a pointer variable)
- Note, however, that this need not be the most efficient approach with respect to run-time (traversing the list means jumping around in memory)
- We will now perform a demo implementation of a linked list for particle tracing.

## Example: Particle Tracing (1/8)

As a first step we define two data structures:

```
// datatype for a single particle
typedef struct party {
    double x, y;
    struct party* next;
} particle_t;

// datatype for all particles (as linked list)
typedef struct {
    particle_t* firstParticle;
    int length;
} particleLine_t;
```

Why do we need `party`?

## Example: Particle Tracing (2/8)

Now a function for generating a new particle:

```
// place a new particle in the flow at given coordinates
particle_t* newParticle( double x, double y ) {

    // allocate new particle
    particle_t* aux = malloc( sizeof( particle_t ) );

    // set particle coordinates
    aux->x = x;
    aux->y = y;

    // no successor, yet
    aux->next = NULL;

    return aux;
}
```

## Example: Particle Tracing (3/8)

And a function for deleting a particle:

```
// delete an existing particle
void delParticle( particle_t* particle ) {
    if( particle != NULL ) {
        free( particle );
    }
}
```

## Example: Particle Tracing (4/8)

```
// The function will insert a particle after predecessor
particle_t* addToList( particle_t* predecessor, particle_t* particle,
                      particleLine_t* pLine ) {

    // first particle needs special treatment
    if( predecessor == NULL ) {
        pLine->firstParticle = particle;
        particle->next = NULL;
    }
    else {
        // get link (address of) successor and set it in the new particle
        particle->next = predecessor->next;

        // insert link to (address of) new particle in predecessor
        predecessor->next = particle;
    }

    // list got longer
    pLine->length++;

    return particle;
}
```

## Example: Particle Tracing (5/8)

```
// remove current particle from list (need predecessor, as we cannot
// go backwards in _singly_ linked list)
void removeFromList( particle_t* predecessor, particle_t* current,
                      particleLine_t* pLine ) {

    // skip over current particle (removing it from the list)
    // if current particle is first in list it needs special treatment
    if( predecessor == NULL ) {
        pLine->firstParticle = current->next;
    }
    else {
        predecessor->next = current->next;
    }

    // delete particle object
    delParticle( current );

    // list got shorter
    pLine->length--;
}
```

## Example: Particle Tracing (6/8)

Working with the list: **list traversal** to output particle info

```
void showParticles( particleLine_t* pLine ) {  
  
    particle_t* current;  
    current = pLine->firstParticle;  
  
    printf( "# -----\n" );  
    printf( "# no. particles in line = %d\n", pLine->length );  
  
    while( current != NULL ) {  
        printf( "particle position: ( %4.2f , %4.2f )\n",  
                current->x, current->y );  
        current = current->next;  
    }  
    printf( "# ----- \n" );  
}
```

Now let us write a test-driver:

```
int main( void ) {

    particle_t* single = NULL;
    particle_t* current = NULL;
    particleLine_t pLine = { NULL, 0 };

    // generate some particles to fill the list initially
    int num = 10;
    current = NULL;
    for( int k = 0; k < num; k++ ) {
        double xPos = (double)rand() / (double)RAND_MAX;
        double yPos = (double)rand() / (double)RAND_MAX;
        single = newParticle( xPos, yPos );
        current = addToList( current, single, &pLine );
    }

    // show what we got
    showParticles( &pLine );
}
```

```
// now remove all particles from list with x > 0.5
current = pLine.firstParticle;
particle_t* previous = NULL;
particle_t* next = NULL;

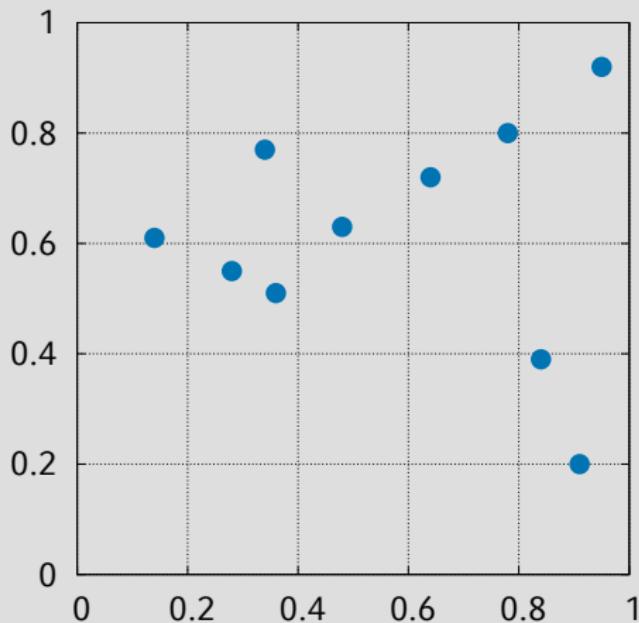
while( current != NULL ) {
    next = current->next;
    if( current->x > 0.5 ) {
        removeFromList( previous, current, &pLine );
    }
    else {
        previous = current;
    }
    current = next;
}
printf( "\n\n" );
showParticles( &pLine );

// add another particle with coords (0.8,0.2) as
// last but one particle
current = pLine.firstParticle;
for( int k = 1; k < pLine.length - 1; k++ ) {
    current = current->next;
}
addToList( current, newParticle( 0.8, 0.2 ), &pLine );
printf( "\n\n" );
showParticles( &pLine );
```

## Example: Particle Tracing (8/8)

we add 10 random points:

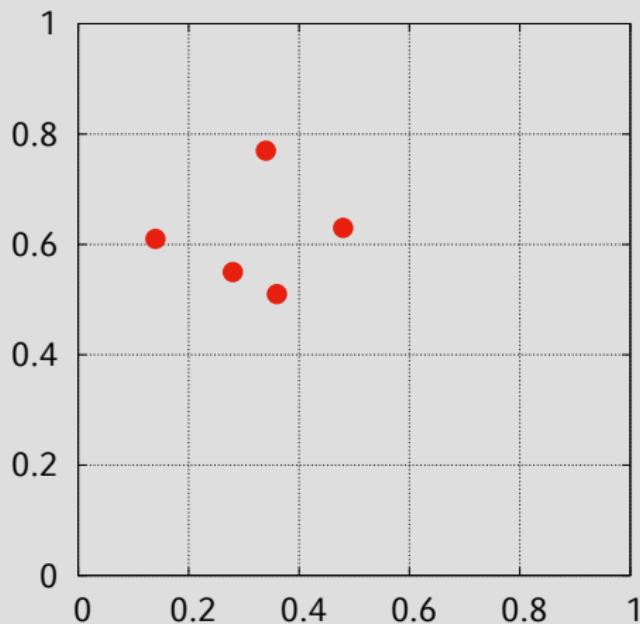
```
# -----
# no. particles in line = 10
particle position: ( 0.84 , 0.39 )
particle position: ( 0.78 , 0.80 )
particle position: ( 0.91 , 0.20 )
particle position: ( 0.34 , 0.77 )
particle position: ( 0.28 , 0.55 )
particle position: ( 0.48 , 0.63 )
particle position: ( 0.36 , 0.51 )
particle position: ( 0.95 , 0.92 )
particle position: ( 0.64 , 0.72 )
particle position: ( 0.14 , 0.61 )
# -----
```



## Example: Particle Tracing (8/8)

now remove all points in right half:

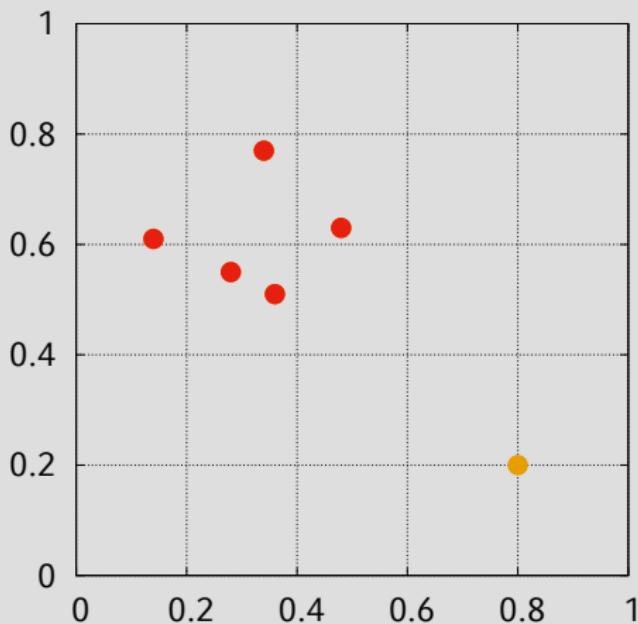
```
# -----
# no. particles in line = 5
particle position: ( 0.34 , 0.77 )
particle position: ( 0.28 , 0.55 )
particle position: ( 0.48 , 0.63 )
particle position: ( 0.36 , 0.51 )
particle position: ( 0.14 , 0.61 )
# -----
```



## Example: Particle Tracing (8/8)

finally add a particle at (0.8, 0.2) as last but one list entry:

```
# -----
# no. particles in line = 6
particle position: ( 0.34 , 0.77 )
particle position: ( 0.28 , 0.55 )
particle position: ( 0.48 , 0.63 )
particle position: ( 0.36 , 0.51 )
particle position: ( 0.80 , 0.20 )
particle position: ( 0.14 , 0.61 )
# -----
```



# Sequence Containers (1/2)

## Container

A container is a class, a data structure, or an abstract data type (ADT) whose instances are collections of other objects.

([Wikipedia](#))

## Sequence

A sequence  $s = \langle e_0, e_1, \dots, e_n \rangle$  is a collection of elements in a specific ordering. Same elements are allowed to appear multiple times at different positions in the sequence.

## Sequence Containers (2/2)

We already know two types of sequence containers:

- singly linked list
  - ▶ allows to store a collection of elements in a specific ordering
  - ▶ does not require to know (maximal) length of sequence a priori
- array
  - ▶ access to element  $e_i$  of sequence  $s$  via `s[i]`
  - ▶ length known a priori: static array
  - ▶ maximal length known a priori: bounded array

Two special types of sequence containers are the **stack** and the **queue**.

## Stacks (1/7)

- example: office clerk
    - ▶ has a stack of documents to process
    - ▶ new documents are put onto top of stack
    - ▶ next document to be processed also taken from top of stack
  - LIFO principle (last in first out)
  - two standard operations:
    - ▶ `push()` = put new element onto top of stack
    - ▶ `pop()` = remove top element from stack

## Stacks (2/7)

- We implement a stack for values of type `elem_t` (float in demo)
- We use an array to implement a `bounded stack` (i.e. stack with a maximal capacity)
- The stack data structure will contain auxilliary management info:

```
typedef struct {
    elem_t* elem;
    int capacity; // capacity of stack
    int top;       // index of top element
} stack_t;
```

- note that we will store the elements on the stack (involves copying), alternatively we could store their addresses (`elem_t**`)

## Stacks (3/7)

- `initStack` will perform memory allocation
- and set the auxilliary info to sensible values
- we define macro `EMPTY_STACK` to be `-1`

```
void initStack( stack_t* stack, int capacity ) {  
    stack->capacity = capacity;  
    stack->top = EMPTY_STACK;  
    stack->elem = malloc( capacity * sizeof( elem_t ) );  
}
```

## Stacks (3/7)

- `initStack` will perform memory allocation
- and set the auxilliary info to sensible values
- we define macro `EMPTY_STACK` to be `-1`  
(as this is no valid value for the top index of the stack)

```
void initStack( stack_t* stack, int capacity ) {  
    stack->capacity = capacity;  
    stack->top = EMPTY_STACK;  
    stack->elem = malloc( capacity * sizeof( elem_t ) );  
}
```

## Stacks (4/7)

- we define two convenience macros to check the state of our stack
- using parameterised macros avoids overhead of external function call

```
#define IS_EMPTY( stack ) (stack->top == EMPTY_STACK)
#define IS_FULL( stack ) (stack->top == stack->capacity - 1)
```

- stack empty, if top element index is EMPTY\_STACK (our convention: -1)
- stack full, if top element index equals (capacity-1)
- in object-oriented terminology `IS_EMPTY()` and `IS_FULL()` are called **getter methods** or **accessor functions**
- they return info on the status of the data-structure (object) but do not alter it.

## Stacks (5/7)

- push places new element (`item`) onto top of stack
- can only push element on stack, if stack is not full!
- need to step top index as part of inserting `item`

```
void push( stack_t* stack, elem_t item ) {  
  
    if ( IS_FULL( stack ) ) {  
        fprintf( stderr, "stack capacity exceeded!\n" );  
        exit( EXIT_FAILURE );  
    }  
  
    stack->top += 1;  
    stack->elem[ stack->top ] = item;  
}
```

## Stacks (6/7)

- pop removes element from top of stack
- value is returned to caller
- can only pop element from stack, if stack is not empty!
- need to decrease top index (no explicit deletion)

```
elem_t pop( stack_t* stack ) {
    if ( IS_EMPTY( stack ) ) {
        fprintf( stderr, "cannot pop from empty stack!\n" );
        exit( EXIT_FAILURE );
    }
    stack->top -= 1;
    return stack->elem[ stack->top + 1 ];
}
```

# Stacks (7/7)

```
stack_t stack = { NULL, 0, 0 };
elem_t aux;

initStack( &stack, 3 );

// operations on stack
push( &stack, 1.0 );
push( &stack, 2.0 );
push( &stack, 3.0 );

aux = pop( &stack );

push( &stack, 4.0 );

aux = pop( &stack );
aux = pop( &stack );
aux = pop( &stack );

clearStack( &stack );
```

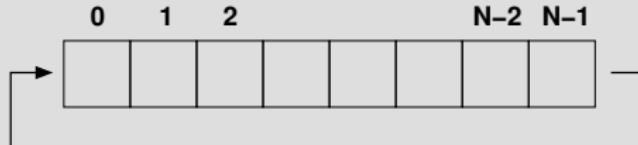
top index	contents of stack
-1	(empty) (before first push)
0	1.0
1	1.0 2.0
2	1.0 2.0 3.0
1	1.0 2.0
2	1.0 2.0 4.0
1	1.0 2.0
0	1.0
-1	(empty) (after last pop)

# Queues

- example: **post office**
  - ▶ customers start queueing at end of line (rear/tail)
  - ▶ first customer in line (front/head) gets served
- **FIFO** principle (first in first out)
- two standard operations:
  - ▶ **enqueue()** = append new element at tail
  - ▶ **dequeue()** = remove element from head
- bounded FIFO queue can be implemented with **circular array**

# Circular Array

- A circular array is just a different interpretation of a standard array.
- We imagine the array being closed, i.e. after the last entry we get to the first again



- Array indices outside  $\mathcal{I} = \{0, 1, \dots, N - 1\}$  can be mapped onto  $\mathcal{I}$  by a modulo operation ( $k \bmod N$ ).

## Queue Example

operation	length	head	tail	queue elements	
enqueue( 1.0 )	1	0	0	1.0 *** *** ***	initial queue empty
enqueue( 2.0 )	2	0	1	1.0 2.0 *** ***	
enqueue( 3.0 )	3	0	2	1.0 2.0 3.0 ***	
enqueue( 4.0 )	4	0	3	1.0 2.0 3.0 4.0	
dequeue()	3	1	3	*** 2.0 3.0 4.0	
dequeue()	2	2	3	*** *** 3.0 4.0	
dequeue()	1	3	3	*** *** *** 4.0	
enqueue( 5.0 )	2	3	0	5.0 *** *** 4.0	wrap-around
enqueue( 6.0 )	3	3	1	5.0 6.0 *** 4.0	
dequeue()	2	0	1	5.0 6.0 *** ***	
dequeue()	1	1	1	*** 6.0 *** ***	
enqueue( 7.0 )	2	1	2	*** 6.0 7.0 ***	
dequeue()	1	2	2	*** *** 7.0 ***	

## Queue Implementation (1/2)

- We store queue elements in an array `queue(0:capacity-1)`. Using 0-based indexing for ease of modulo operation.
- We employ this queue data structure with management info

```
typedef struct {
    elem_t* elem;
    int capacity;    // capacity of queue
    int head;        // index of head of queue
    int tail;        // index of tail of queue
    int length;      // current length of queue
} queue_t;
```

- Current queue elements reside in `elem[head] ↔ elem[tail]`, including possibly a wrap-around.

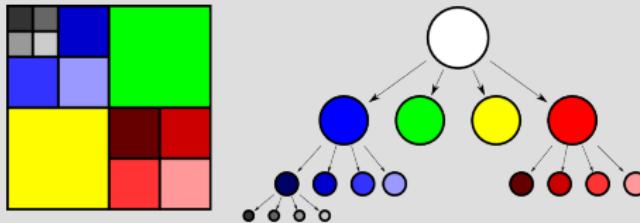
## Queue Implementation (2/2)

- as an example we consider `enqueue()` routine.
- `item` contains new value to be added **after tail**
- for this we increment tail and perform a modulo operation on it:

$$k \bmod q_{\text{cap}} \in \{0, 1, \dots, q_{\text{cap}} - 1\}$$

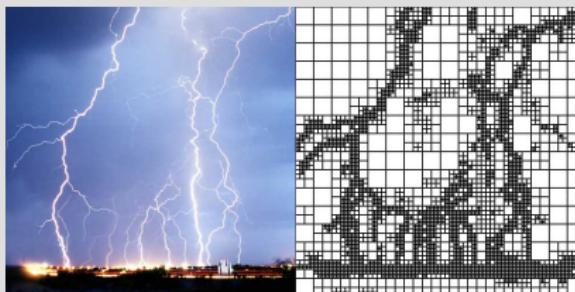
```
void enqueue( queue_t* queue, elem_t new ) {
    if( queue->length < queue->capacity ) {
        queue->tail = (queue->tail + 1) % queue->capacity;
        queue->elem[ queue->tail ] = new;
        queue->length += 1;
    }
    else {
        fprintf( stderr, "queue capacity exceeded!\n" );
        exit( EXIT_FAILURE );
    }
}
```

# Quadtrees



- a quadtree is a specialised form of tree structure
- each node of the tree is either a leaf (no children) or has four children
- identical concept for 3D is an octree
- used e.g. in image processing, area indexing (GIS), collision detection, hidden surface removal of terrain data, mesh generation

# QuadTrees — Examples

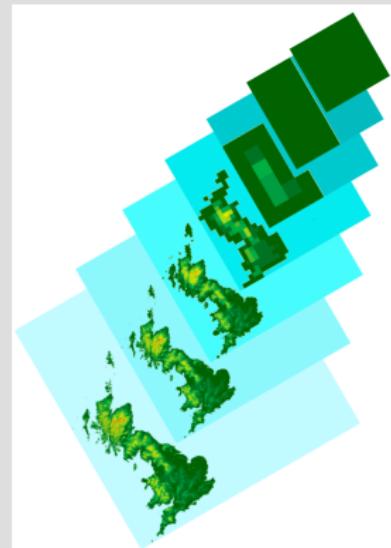


Originalbild

Aus Konturwerten erstellter Quadtree  
(Größe: 1229 Byte)

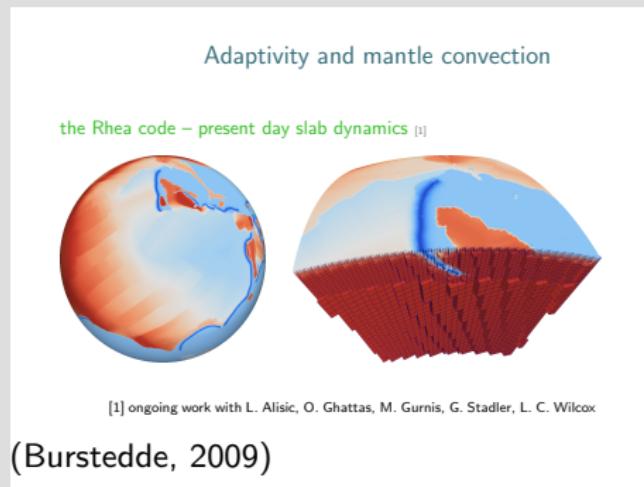
Das Bild zeigt einen aus den Konturwerten berechneten Quadtree. Zu beachten ist, dass bei starken Konturen (hier: die Blitze) die Quadrate sehr klein sein müssen, was sich im Speicherverbrauch niederschlägt.

Lightning



Great Britain  
Kidner et al.

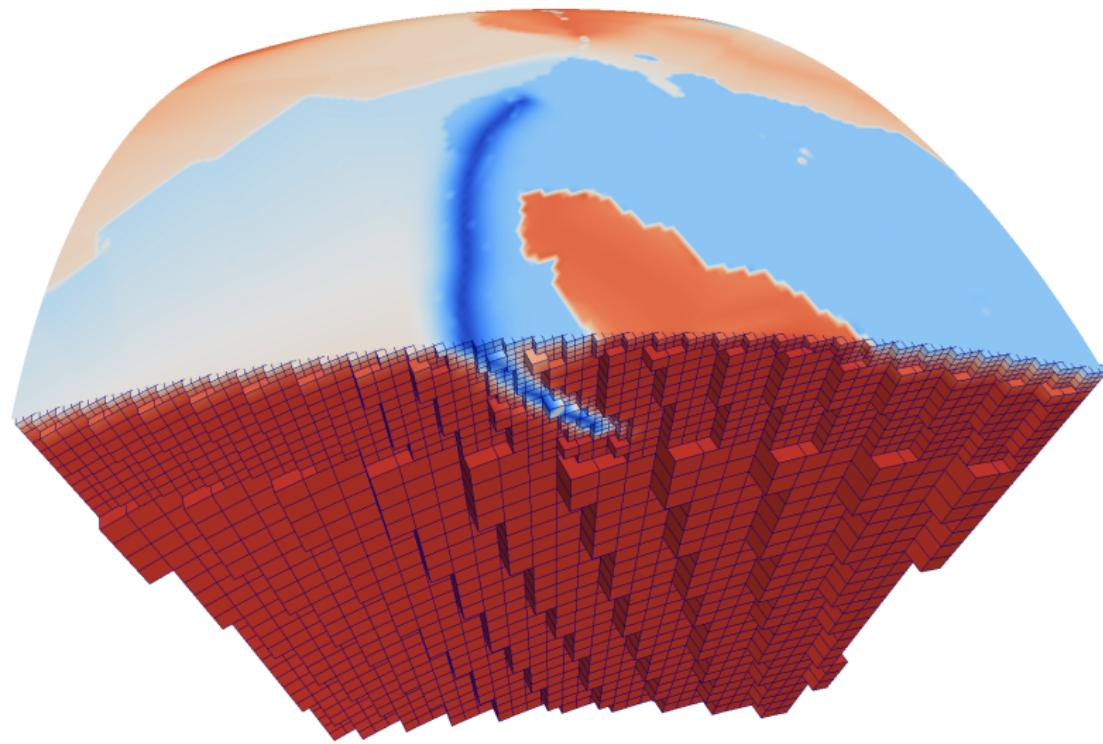
# Octrees — Example



- Octrees are the 3D extension of quadtrees
- are used in many applications such as
  - ▶ spatial indexing
  - ▶ efficient collision detection in three dimensions
  - ▶ view frustum culling
  - ▶ fast Multipole Method
  - ▶ unstructured grid
  - ▶ finite element analysis

## Octrees — Example

### slab dynamics [1]



# Dense versus Sparse Matrices

## (Weak) Definition:

A matrix is called sparse when the number of non-zero matrix entries is significantly smaller than the number of zero matrix entries.

On the algorithmic side sparse matrices pose completely different challenges than dense matrices.

# Dense versus Sparse Matrices

## (Weak) Definition:

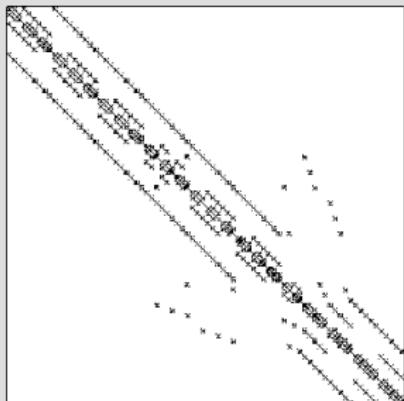
A matrix is called sparse when the number of non-zero matrix entries is significantly smaller than the number of zero matrix entries.

On the algorithmic side sparse matrices pose completely different challenges than dense matrices. E.g.

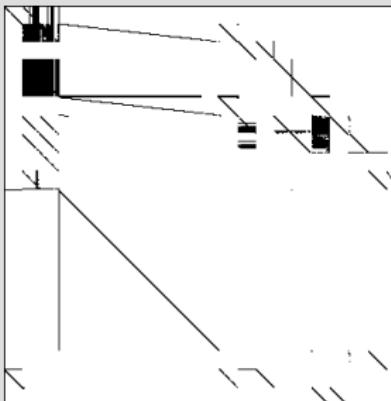
- Storing a sparse matrix as 2D-array is not economic (due to the large number of zeros) and in most cases not even feasible
- Fast methods for solving  $Ax = b$  must take the sparsity into account, i.e. you only want to deal with the non-zero entries

# Examples from the MatrixMarket

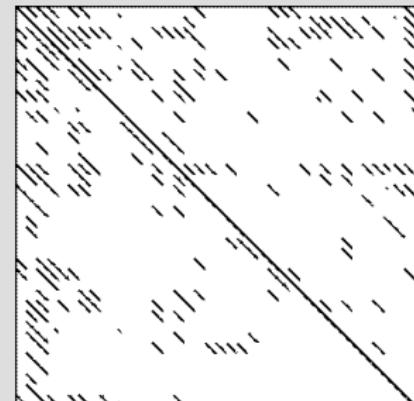
<http://math.nist.gov/MatrixMarket/>



Structural  
Engineering  
(NOS5)



Economic  
Modelling  
(ORANI678)



Weather  
Simulation  
(FS 760 3)

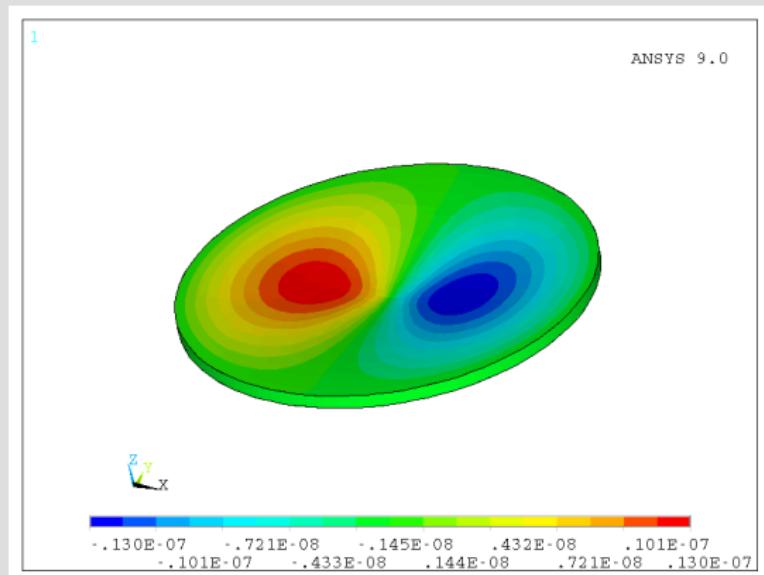
# Where do Sparse Matrices come from?

A major source of sparse matrices is the discretisation of **partial differential equations**.

In practice the four most important discretisation methods are:

Finite Differences	<b>sparse</b> matrix	}	volume discretisation techniques
Finite Elements	<b>sparse</b> matrix		
Finite Volumes	<b>sparse</b> matrix		
Boundary Elements	<b>dense</b> matrix	}	surface discretisation technique

# Example: Structural Mechanics



$u_x = u_y = u_z = 0$   
on boundary

load in centre

standard discretisation  
with tri-linear Finite  
elements on hexaeders

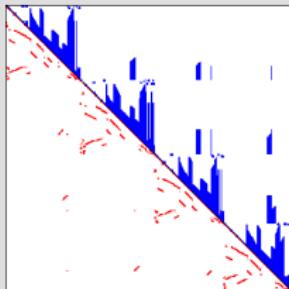
# nodes	3,726
# elements	2,370
<i>n</i>	10,170
nnz	619,038
sparsity	0.60%

## Example: Re-ordering

Cholesky decomposition  $A = LL^T$  with  $\dim(A) = 10,170$   
non-zeros to store  $A = 314,604$

## Example: Re-ordering

Cholesky decomposition  $A = LL^T$  with  $\dim(A) = 10,170$   
non-zeros to store  $A = 314,604$



'original'

non-zeros in  $L$

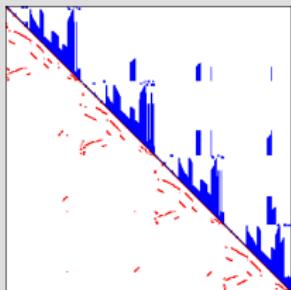
5,749,983

increase by

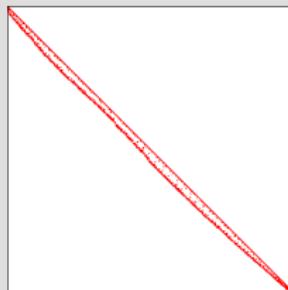
18.3

## Example: Re-ordering

Cholesky decomposition  $A = LL^T$  with  $\dim(A) = 10,170$   
non-zeros to store  $A = 314,604$



'original'



Cuthill-McKee

non-zeros in L

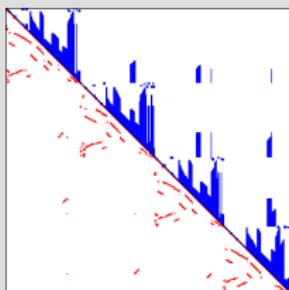
5,749,983

increase by

18.3

## Example: Re-ordering

Cholesky decomposition  $A = LL^T$  with  $\dim(A) = 10,170$   
non-zeros to store  $A = 314,604$



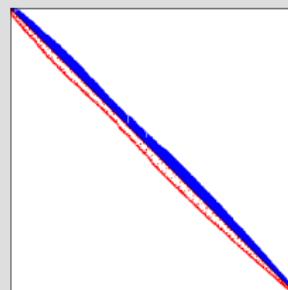
'original'

non-zeros in L

5,749,983

increase by

18.3



Cuthill-McKee

non-zeros in L

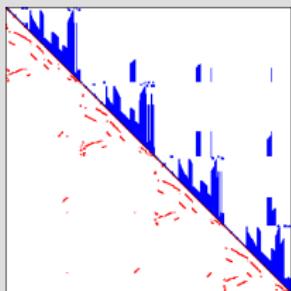
3,581,856

increase by

11.4

## Example: Re-ordering

Cholesky decomposition  $A = LL^T$  with  $\dim(A) = 10,170$   
non-zeros to store  $A = 314,604$



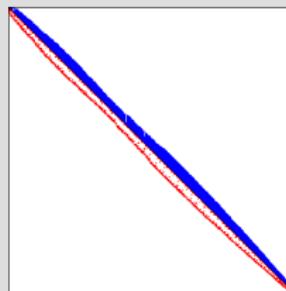
'original'

non-zeros in L

5,749,983

increase by

18.3



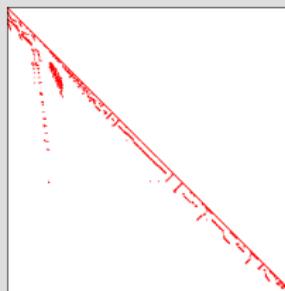
Cuthill-McKee

non-zeros in L

3,581,856

increase by

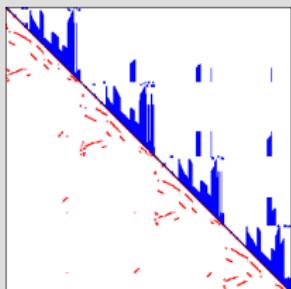
11.4



Sloan

## Example: Re-ordering

Cholesky decomposition  $A = LL^T$  with  $\dim(A) = 10,170$   
non-zeros to store  $A = 314,604$



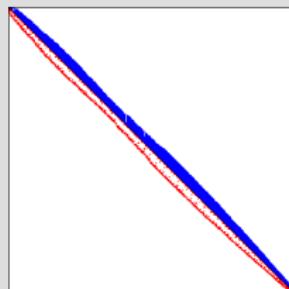
'original'

non-zeros in L

5,749,983

increase by

18.3



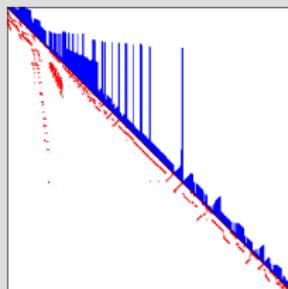
Cuthill-McKee

non-zeros in L

3,581,856

increase by

11.4



Sloan

non-zeros in L

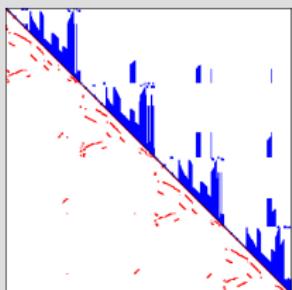
4,133,448

increase by

13.1

## Example: Re-ordering

Cholesky decomposition  $A = LL^T$  with  $\dim(A) = 10,170$   
non-zeros to store  $A = 314,604$



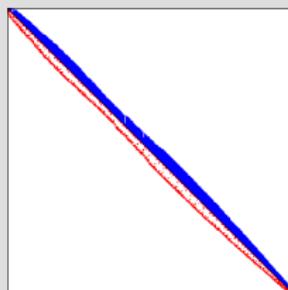
'original'

non-zeros in L

5,749,983

increase by

18.3



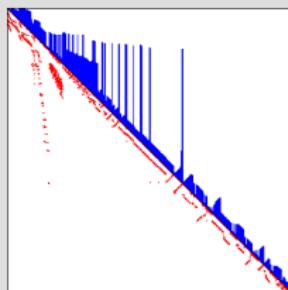
Cuthill-McKee

non-zeros in L

3,581,856

increase by

11.4



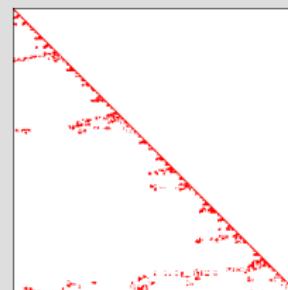
Sloan

non-zeros in L

4,133,448

increase by

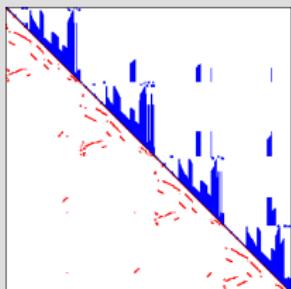
13.1



minimum degree

## Example: Re-ordering

Cholesky decomposition  $A = LL^T$  with  $\dim(A) = 10,170$   
non-zeros to store  $A = 314,604$



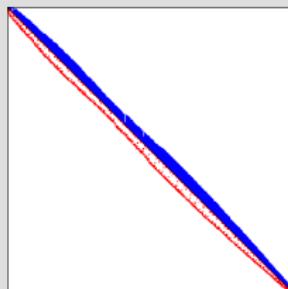
'original'

non-zeros in L

5,749,983

increase by

18.3



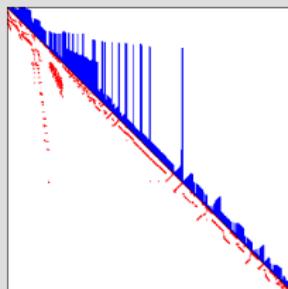
Cuthill-McKee

non-zeros in L

3,581,856

increase by

11.4



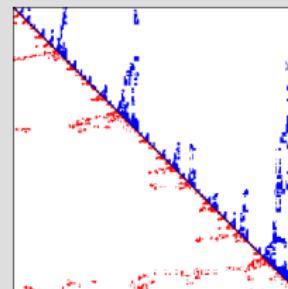
Sloan

non-zeros in L

4,133,448

increase by

13.1



minimum degree

non-zeros in L

2,029,509

increase by

6.5

# Practical Issues

The question **how to store a sparse matrix** turns out to be non-trivial.

- storing it as a 2D array ← **impossible**
- answer strongly depends on what you want to do with it

## Approach I: Index scheme

For each non-zero entry  $a_{ij}$  store the triplet:

row index $i$	column index $j$	non-zero value $a_{ij}$
---------------	------------------	-------------------------

- we need three arrays of length nnz (number of non-zeros)
- each triplet requires (4+4+8) byte per entry (on typical 32-bit architecture)
- total storage costs:  $16 \cdot \text{nnz}$  bytes

# Index scheme (cont.)

## Matrix from Structural Mechanics example

```
%%MatrixMarket matrix coordinate real symmetric
%
% Matrix exported by OLAS
%
10170 10170 314604
 1      1  8.2990067682228677e+04
 1      2 -4.9696384386100754e+03
 1      3 -2.7606714995454240e-01
 1      4  9.7198803223981122e+03
 1      5 -6.1171895993969656e+03
 1      6 -2.7255939150670642e+03
 1      7 -1.7210872534276655e+03
 1      8  2.5811602286800858e+03
 1      9  8.7842380729450872e+02
```

problem with index scheme is that it does not give any structural support, e.g. no indication where a row starts

(think of matrix-vector product  $A \cdot u$ )

we can also do a little better on storage

requires 9.4 MB of storage

## Approach II: Compressed Row Storage

$$\begin{bmatrix} -4 & 8 & 7 & \cdot \\ 5 & -6 & \cdot & 2 \\ 3 & \cdot & -1 & \cdot \\ \cdot & 10 & \cdot & -9 \end{bmatrix}$$

CRS format employs three arrays:

1 2 3 4 5 6 7 8 9 10

val = 

-4	8	7	5	-6	2	3	-1	10	-9
----	---	---	---	----	---	---	----	----	----

cidx = 

1	2	3	1	2	4	1	3	2	4
---	---	---	---	---	---	---	---	---	---

rptr = 

1	4	7	9	11
---	---	---	---	----

## Compressed Row Storage (cont.)

The idea of the compressed row storage format is to:

- to store for all non-zero entries their **value and column index**
- entries for a row are stored contiguously (but in no special order)
- rows are stored one after the other lexicographically
- for each row we have a 'pointer' where the row starts in the other two arrays

## Compressed Row Storage (cont.)

- CRS requires a little less memory:

array	contents	type	length
val	values	(double)	nnz
cidx	column indices	(integer)	nnz
rptr	row start indices	(integer)	$n + 1$

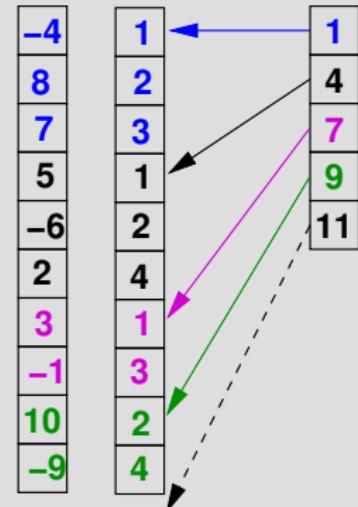
$$\rightarrow (8 + 4)nnz + 4(n + 1) \quad (\rightarrow 7.1 \text{ MB in example})$$

- CRS provides **structural information**
  - ▶ we know where a row starts
  - ▶ we know how many non-zeros are in row  $k$  from `rptr[k+1] - rptr[k]`

# Matrix-Vector Product with CRS

The product  $Au = v$  can be implemented as

```
for i = 1 to n do
    sum = 0
    k = rowPtr[i]
    rs = rowPtr[i+1] - k
    for j = 1 to rs do
        sum = sum + val[k] · u[cidx[k]]
        k = k + 1
    end for
    v[i] = sum
end for
```



## Other Sparse Storage Formats

There is a multitude of different storage formats for sparse matrices

- CRS sub-variants (like storing the diagonal entry in first position for each row)
- Compressed Column Storage (CCS)
- Compressed Diagonal Storage (CDS)
- Jagged Diagonal Storage (JDS)
- Skyline Storage

which all have their pros and cons.

# Resources

- Bebis, Univ. of Nevada, [CS308: Data Structures](#)
- Morris, Univ. of Auckland, [PLDS210: Programming Languages and Data Structures](#)
- [Wikipedia](#) entry on *Data Structure* and associated entries
- Barrett, Berry, Chan et al., '[Templates for the Solution of Linear Systems: Building Blocks for Iterative Methods](#)', SIAM, 1994 ([www.netlib.org](http://www.netlib.org))