

P6 – Scientific Programming

Marcus Mohr
Jens Oeser

Geophysics Section
Department of Earth and Environmental Sciences
Ludwig-Maximilians-Universität München

SoSe 2021

Part #16

Introduction to GNU Make

Automated Builds, Targets,
Dependencies and Macros

Automated Builds

- Problems:
 - ▶ I don't want to type `ifort -align dcommons -DBoussinesq -I/usr/lib/mpich-intel/include -O3 -xW -axW -c driver.f` each time I change something in `driver.f` in order to re-compile.
 - ▶ I don't want to re-compile files that did not change to get new executable (but which did not change?)
 - ▶ Each time a new seismogram is added to my database `myReport.tex` should be updated using the `evalScript.sh`.
- One possible solution is to use automated build systems.
- We today take a look at **GNU Make**

Acknowledgements



- 95% of the following slides are based on the course on [Software Carpentry](#) developed by Brent Gorda and Greg Wilson
- The course is publicly available at <http://swc.scipy.org/>
- and was supported by the Python Software Foundation and the University of Toronto
- Should take a look at the material before your master's thesis!

License

applicable to slides of part 16 of the course

Copyright (c) 2005 Python Software Foundation

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

Introduction

- Most languages require you to compile programs before running them
 - ▶ Typing `gcc -c -Wall -ansi -I/pkg/chempak/include dat2csv.c` once is bad enough.
 - ▶ Typing it dozens of times as you edit and debug is **tedious** and **error-prone**.
- Most large programs contain dependencies
 - ▶ Module A uses modules B and C, B uses D and E, C uses E and F, etc.
 - ▶ If E changes, ought to recompile B and C, then A.

Rule: Anything worth repeating is worth automating!

- A standard way and place to save project-related commands ...
- ... that keeps track of what depends on what

You may fall asleep now, if ...

- You know what a Makefile is
- You know how to write a rule
- You know how dependencies affect the order of command execution
- You know how to define macros
- You know how to use automatic variables
- You know how to write a generic rule

Automate, Automate, Automate

- Tools that manage repetitive tasks and their dependencies are usually called **build tools**
 - ▶ Originally developed to rebuild software packages
 - ▶ Can equally well be used to update web site content, run backups, ...
- Such a tool must have:
 - ▶ A way to describe what things to do
 - ▶ A way to specify the dependencies between them

Make

- Most widely used build tool is Make
 - ▶ Invented at Bell Labs in 1975 by Stuart Feldman.
 - ▶ He went on to become a vice-president at IBM, which shows you how far a good tool can take you.
- The good news: Make is freely available for every major platform, and very well documented.
- The bad news is Make's syntax
 - ▶ Over 30 years, it has grown into a little programming language
 - ▶ We will ignore advanced features for now
 - ▶ For alternatives/improvements see e.g. [CMake](#), [SCons](#) or [Apache Ant](#) or integrated development environments (IDEs) such as [eclipse](#).

Our Example

- Running example: Nigel is studying organic fullerene production
 - Automated laboratory equipment runs experiments in batches to create files like this:


```
Time: 1.2271 Concentration: 0.0050 Yield: 11.41
Time: 2.5094 Concentration: 0.0055 Yield: 11.20
Time: 3.7440 Concentration: 0.0060 Yield: 10.90
```
- Each experiment produces 20-30 files
 - Use comma-separated values (CSV) for the table
 - May eventually have several thousand of them
- Want to:
 - Generate tables showing the results for particular trials using a program called `dat2csv`
 - Update a file showing the correlation between concentrations and yields based on those tables

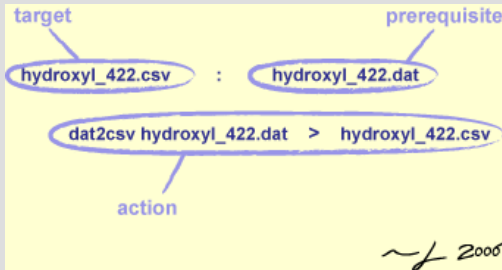
Hello, Make

- Put the following into a Makefile called **hello.mk**:

```
hydroxyl_422.csv : hydroxyl_422.dat
    dat2csv hydroxyl_422.dat > hydroxyl_422.csv
```
- Must indent with a **tab character**: not eight spaces, or a mix of spaces and tabs
 - ▶ Yes, it's a wart, but we're stuck with it
- Run **make -f hello.mk**
 - ▶ Make sees that the CSV file depends on the data file
 - ▶ Since the CSV file doesn't exist, Make runs

```
dat2csv hydroxyl_422.dat > hydroxyl_422.csv
```
- Run **make -f hello.mk** again
 - ▶ **hydroxyl_422.csv** is newer than **hydroxyl_422.dat**, Make does not run the command again

Terminology



- `hydroxyl_422.csv` is the **target** of the rule
- `hydroxyl_422.dat` is its **prerequisite**

- The compilation command is the rule's **action**
Make runs it on your behalf, just as the shell runs the commands you type.

Multiple Targets

- Makefiles usually contain multiple rules

```
hydroxyl_422.csv : hydroxyl_422.dat  
                  dat2csv hydroxyl_422.dat > hydroxyl_422.csv
```

```
methyl_422.csv : methyl_422.dat  
                dat2csv methyl_422.dat > methyl_422.csv
```

- When you run `make -f double.mk`, only `hydroxyl_422.csv` is compiled
 - ▶ The first rule in the Makefile specifies the **default target**
 - ▶ Unless you tell it otherwise, that's all Make will update
- Have to run `make -f double.mk methyl_422.csv` to build `methyl_422.csv`

Phony Targets

- Running Make separately for each target would hardly count as “automation”
- Solution: define a **phony target** that:
 - ▶ Depends on all the things you want to recompile, but doesn't correspond to any files
 - ▶ It can never be up to date, so making it will always executes its actions

```
all : hydroxyl_422.csv methyl_422.csv
```

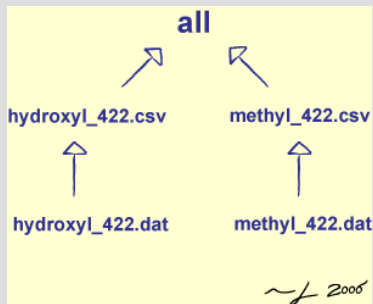
```
hydroxyl_422.csv : hydroxyl_422.dat
                   dat2csv hydroxyl_422.dat > hydroxyl_422.csv
```

```
methyl_422.csv : methyl_422.dat
                 dat2csv methyl_422.dat > methyl_422.csv
```

- `make -f phony.mk all` now creates both .csv files

Dependencies

- Note how one target can depend on others
 - ▶ `all` depends on `hydroxyl_422.csv` and `methyl_422.csv`
 - ▶ Each of these depends on (i.e., must be newer than) the corresponding `.dat` file
- Can visualize dependencies as a directed graph
 - ▶ Each file is represented by a node
 - ▶ Dependencies are then the graph's arcs/edges



Updating Dependencies

- Make's built-in processing cycle:
 - ▶ Follow links top-down to find direct and indirect dependencies
 - ▶ Execute actions bottom-up to update
- Make can execute actions in any order it wants to, as long as it doesn't violate dependency ordering
 - ▶ Could update either `hydroxyl_422.csv` or `methyl_422.csv` first
 - ▶ But has to update both **before** "updating" `all`

Conventions

- If you run `make` with no arguments, it automatically looks for a file called `Makefile`
 - ▶ So most projects use that name for their Makefile
 - ▶ And remember, without an explicit target name, `make` only updates the first one it finds
- Typical phony targets in a typical Makefile include:
 - ▶ `all`: recompile everything
 - ▶ `clean`: delete all temporary files, and everything produced by compilation
 - ▶ `install`: copy files to system directories
- Many open source packages can be installed by typing:
 - ▶ `make configure` (more often `configure`)
 - ▶ `make`
 - ▶ `make test`
 - ▶ `make install`

Automatic Variables

Make defines automatic variables to represent parts of rules

- Values re-set for each rule
- Unfortunately, names are very cryptic

\$@	the rule's target
\$<	the rule's first prerequisite
\$?	all of the rule's out-of-date prerequisites
\$^	all prerequisites

Automatic Variables Example

- Rewrite the Makefile using automatic variables

```
all : hydroxyl_422.csv methyl_422.csv
```

```
hydroxyl_422.csv : hydroxyl_422.dat
```

```
@dat2csv $< > $@
```

```
methyl_422.csv : methyl_422.dat
```

```
@dat2csv $< > $@
```

```
clean :
```

```
@rm -f *.csv
```

- By default, Make echoes actions before executing them
 - ▶ Putting @ at the start of the action line prevents this
- And add a phony target clean to tidy up generated files
 - ▶ Question: why `rm -f` instead of just `rm`?

Pattern Rules

- Most files of similar type in a project are processed the same way
 - ▶ e.g., typically compile all Fortran or Java files with the same options
- Write a pattern rule to describe the general case

```
all : hydroxyl_422.csv methyl_422.csv
```

```
%.csv : %.dat
```

```
@dat2csv $< > $@
```

- The **wildcard %** represents the stem of the file's name in the target and prerequisites
- Must use automatic variables in the actions
 - This is why they were invented
- Pattern rules are **implicit rules** defined by the user; make also has predefined implicit rules

Adding More Dependencies (1/2)

Now create a summary for each set of experiments:

- Use **summarize** command to combine data from `hydroxyl_422.csv` and `hydroxyl_480.csv`
- Output is `hydroxyl_all.csv`
- Perform same calculation for methyl files

Adding More Dependencies (2/2)

- Updated Makefile is a simple extension of what we've seen before:

```
all : hydroxyl_all.csv methyl_all.csv

%_all.csv : %_422.csv %_480.csv
    summarize $^ > $@

%.csv : %.dat
    dat2csv $< > $@

clean :
    @rm -f *.csv
```

- The rule for `%_all.csv` takes precedence over the rule for `%.csv`.
- Make uses the most specific rule available.

Tidying Up

- What happens when this file is executed for the first time?

```
$ make -f depend.mk
dat2csv hydroxyl_422.dat > hydroxyl_422.csv
dat2csv hydroxyl_480.dat > hydroxyl_480.csv
summarize hydroxyl_422.csv hydroxyl_480.csv > hydroxyl_all.csv
dat2csv methyl_422.dat > methyl_422.csv
dat2csv methyl_480.dat > methyl_480.csv
summarize methyl_422.csv methyl_480.csv > methyl_all.csv
rm hydroxyl_480.csv methyl_422.csv hydroxyl_422.csv methyl_480.csv
```

- Make automatically removes intermediate files created by pattern rules when it's done
- Question: how do you prevent this?

Tidying Up

- What happens when this file is executed for the first time?

```
$ make -f depend.mk
dat2csv hydroxyl_422.dat > hydroxyl_422.csv
dat2csv hydroxyl_480.dat > hydroxyl_480.csv
summarize hydroxyl_422.csv hydroxyl_480.csv > hydroxyl_all.csv
dat2csv methyl_422.dat > methyl_422.csv
dat2csv methyl_480.dat > methyl_480.csv
summarize methyl_422.csv methyl_480.csv > methyl_all.csv
rm hydroxyl_480.csv methyl_422.csv hydroxyl_422.csv methyl_480.csv
```

- Make automatically removes intermediate files created by pattern rules when it's done
- Question: how do you prevent this? make 'em **.PRECIOUS : %.csv**

Defining Macros (1/2)

- Often we want to define variables inside a Makefile
 - ▶ The output directory, the optimization flags for the compiler, ...
 - ▶ Experience shows that

Rule: Anything repeated in two or more places
will eventually be wrong in at least one

- Solution: define variables (usually called **macros**)
 - ▶ Remember: Make is a little programming language
 - ▶ Change behaviour by changing one value in one place

Defining Macros (2/2)

```
INPUT_DIR = /lab/gamma2100
OUTPUT_DIR = /tmp
```

```
all : ${OUTPUT_DIR}/hydroxyl_all.csv ${OUTPUT_DIR}/methyl_all.csv
```

```
${OUTPUT_DIR}/%_all.csv : ${OUTPUT_DIR}/%_422.csv ${OUTPUT_DIR}/%_480.csv
    @summarize $^ > $@
```

```
${OUTPUT_DIR}/%.csv : ${INPUT_DIR}/%.dat
    @dat2csv $< > $@
```

- To get value, put a **\$** in front of the name and parentheses or braces around it; You can use **\$(XYZ)** or **\${XYZ}**
- Without the parentheses Make interprets **\$XYZ** as the value of **X**, followed by the characters **YZ**. *Yes, it's another wart.*

Passing Values to Make (1/2)

- Sometimes it is useful to pass values into Make when invoking it
 - ▶ e.g. change the input directory
- Instead of editing the Makefile, specify `name=value` pairs on the command line
 - ▶ define a macro with the default value
 - ▶ override it when you want to
- Thus
 - ▶ `make -f macro.mk` sets `INPUT_DIR` to `/lab/gamma2100`
 - ▶ but `make INPUT_DIR=/newlab -f macro.mk` uses `/newlab`

Passing Values to Make (2/2)

```
VAL = original
```

```
echo :
```

```
    @echo "VAL is" ${VAL}
```

```
==> make -f env.mk
```

```
VAL is original
```

```
==> make VAL=changed -f env.mk
```

```
VAL is changed
```

- Make also examines environment variables
- You can refer to `${HOME}` in a Makefile without having defined it

Precedence

command line > value set in makefile > environment variable

GNU Make

- <http://www.gnu.org/software/make/>
- GNU make is the Make version of the GNU project
- It is frequently used in conjunction with the GNU build system
- Ubiquitous on Linux and most HPC systems
- Its departures from traditional make are most noticeable in
 - ▶ pattern-matching in dependency graphs and build targets
 - ▶ as well as a number of functions which may be invoked to have the make utility do things like collect a list of all files in the current directory, ...

Function Example

Turn **hydroxyl** into `/tmp/hydroxyl_all.csv` and **methyl** into `/tmp/methyl_all.csv`

```
INPUT_DIR = /lab/gamma2100
OUTPUT_DIR = /tmp
CHEMICALS = hydroxyl methyl
SUMMARIES = $(addprefix ${OUTPUT_DIR}/,$(addsuffix _all.csv,$${CHEMICALS}))

all : ${SUMMARIES}

${OUTPUT_DIR}/%_all.csv : ${OUTPUT_DIR}/%_422.csv ${OUTPUT_DIR}/%_480.csv
    @summarize $^ > $@

${OUTPUT_DIR}/%.csv : ${INPUT_DIR}/%.dat
    @dat2csv $< > $@

clean :
    @rm -f *.csv
```

Commonly-Used Functions

Function	Purpose
<code>\$(addprefix prefix,filenames)</code>	add a prefix to each filename in a list
<code>\$(addsuffix suffix,filenames)</code>	add a suffix to each filename in a list
<code>\$(dir filenames)</code>	extract directory part of each filename in a list
<code>\$(filter pattern,text)</code>	Keep words in text that match pattern
<code>\$(filter-out pattern,text)</code>	Keep words in text that don't match pattern
<code>\$(sort text)</code>	Sort the words in text, removing duplicates
<code>\$(strip text)</code>	Remove leading and trailing whitespace from text
<code>\$(subst from,to,text)</code>	Replace from with to in text
<code>\$(wildcard pattern)</code>	Create a list of filenames that match a pattern

Pros and Cons

- Pro

- ▶ Simple things are simple to do ...
- ▶ ...and not too difficult to read ...
- ▶ ...especially compared to the alternatives

- Con

- ▶ The syntax is unpleasant
- ▶ Complex things are difficult to read ...
- ▶ ...and even more difficult to debug
 - ▶ Best you can do is use echo to print things as Make executes
- ▶ Not really very portable
 - ▶ Hands commands to the shell for execution
 - ▶ But commands use different flags on different operating systems
 - ▶ Do you use del or rm to delete files?

Summary

- Two rules for healthy software projects:
 - ▶ Every repetitive task is done through the build system
 - ▶ Never commit anything to version control repository that breaks the build
- Remember: a Makefile is a program
 - ▶ So give your build the same careful attention you'd give any other programming problem

- The traditional GNU Build Systems is composed of two components:

where autoconf takes care of adapting makefiles to the target environment (OS, hardware & compiler).

- | | |
|--------------|---|
| SCons : | config in Python, replaces both make + autoconf with one tool |
| CMake : | alternative to autoconf; needs an underlying build system (such as make); will automatically generate makefiles |
| ninja : | a make replacement; not intended for manual use; works nicely together with CMake |
| Apache Ant : | (Another Neat Tool); make replacement; uses XML and Java |

Obtaining & Building HyTeG

- General instructions:

Click on link: <https://i10git.cs.fau.de/hyteg/hyteg>

- Preparations on our system:

- 1 `git clone --recurse-submodules
https://i10git.cs.fau.de/hyteg/hyteg.git`
- 2 `mkdir hyteg-build`
- 3 `cd hyteg-build`
- 4 `module load gcc/8.3.0 mpi.omp-gcc/8.3.0 petsc/3.11.0`
- 5 `run cmake`
- 6 `run make (or ninja)`

Configuring the Build

HyTeG uses CMake as build system. CMake is an open-source, cross-platform family of tools designed to build, test and package software (by Kitware). Needs an underlying native build environment (default: GNU Make).



```
cmake -DCMAKE_CXX_COMPILER=g++  
      -DCMAKE_CXX_COMPILER_LAUNCHER=ccache  
      -GNinja  
      -DHYTEG_BUILD_WITH_EIGEN=yes  
      -DHYTEG_BUILD_WITH_PETSC=yes  
      -DCMAKE_BUILD_TYPE=Debug  
      ../HyTeG
```

Executing the Build (make)

We build a single app from scratch (involves building libhyteg.a):

```
cd apps; make -j 10 PolarLaplacian
```

```
Scanning dependencies of target core
Scanning dependencies of target hyteg
[ 1%] Building CXX object walberla/src/core/CMakeFiles/core.dir/Abort.cpp.o
[ 1%] Building CXX object walberla/src/core/CMakeFiles/core.dir/Conversion.cpp.o
[ 1%] Building CXX object walberla/src/core/CMakeFiles/core.dir/Environment.cpp.o
[ 1%] Building CXX object walberla/src/core/CMakeFiles/core.dir/DataTypes.cpp.o
[ 1%] Building CXX object walberla/src/core/CMakeFiles/core.dir/GetPID.cpp.o
[ 1%] Building CXX object walberla/src/core/CMakeFiles/core.dir/RandomUUID.cpp.o
:
:
[100%] Building CXX object src/hyteg/CMakeFiles/hyteg.dir/primitivestorage/loadbalancing/SimpleBalancer.cpp.o
[100%] Linking CXX static library libhyteg.a
[100%] Built target hyteg
Scanning dependencies of target PolarLaplacian
[100%] Building CXX object apps/CMakeFiles/PolarLaplacian.dir/PolarLaplacian.cpp.o
[100%] Linking CXX executable PolarLaplacian
```

Executing the Build (ninja)

We build a single app from scratch (involves building libhyteg.a):

```
ninja -j 10 apps/PolarLaplacian
```

```
⋮
```

```
[8/1417] Building CXX object walberla/...les/core.dir/debug/TestSubsystem.cpp.o
```

```
⋮
```

```
[1182/1417] Building CXX object ...generatedKernels/sor_3D_macroface_P2_update_vertexdofs_impl_023_310.cpp.o
```

```
⋮
```

```
[1414/1417] Linking CXX static library src/hyteg/libhyteg.a
```

```
[1415/1417] Building CXX object apps/CMakeFiles/PolarLaplacian.dir/PolarLaplacian.cpp.o
```

```
[1416/1417] Linking CXX executable apps/PolarLaplacian
```

